

Maps and Sorted Maps; Binary Search Trees



Map ADT

A MAP is an ADT to efficiently store and retrieve values based on a uniquely identifying **search key**.

It stores key-value pairs (k,v) , which we call **entries**.

Keys are unique/no repeats (they uniquely identify the value); a key is mapped to a value.

The main operations of a MAP are **searching**, **inserting**, and **deleting** items.

Examples: student records, user accounts, etc

Typical keys are username, user ID, etc.

Maps are also known as **associative arrays**.

Dictionary ADT is related, although it normally refers to a similar ADT that allows repeated keys.

The MAP ADT methods:

`get(k)`: returns the value v associated to key k , if such entry exists; otherwise returns null.

`put(k, v)`: if M does not have an entry with key k , then adds (k, v) and returns null; otherwise it replaces with v the value of the entry with key equal to k and returns the old value.

`remove(k)`: removes from M the entry with key k and returns its value; if M has no such entry, then returns null.

`size()`: returns the number of entries in M .

`isEmpty()`: boolean indicating if M is empty.

`keySet()`, `values()`, `entrySet()` returns an iterable collection of keys, values, key-value entries (respectively) stored in M .

MAP ADT: examples

Applications/examples:

- **University information system:**

key= student id

value= student record (name, address, course grades)

- A **domain name system (DNS)** maps a host name (key, e.g. www.wiley.com) to a IP address (value, e.g. 208.215.179.146)

- A **social media site** maps a username which is the key (usually nonnumeric) to the user info which is the value (typically tons of personal info)

SORTED MAP ADT methods:

In addition to the MAP methods:

`get(k); put(k, v); remove(k); size(); isEmpty();
keySet(); values(); entrySet()`

A SORTED MAP also provides:

`firstEntry(), lastEntry()`: returns the entry with smallest key, largest key (respectively), or null if the map is empty.

`subMap(k1,k2)`: returns an iterable list with all the entries greater than or equal to `k1`, but strictly less than `k2`.

`lowerEntry(k), higherEntry(k), floorEntry(k), ceilingEntry(k)`
return the entry with, respectively:
the greatest key $< k$, the smallest key $> k$,
the greatest key $\leq k$, the smallest key $\geq k$.

Implementing MAPs:

- Using an Unordered Sequence
- Using an Ordered Sequence
- Using Search Trees - binary search trees, AVL trees, red-black trees, (2,4)-trees
(starts this lecture and continue on next ones)
- Using Hash Table
(discussed at a later lecture)

Implementing SORTED MAPs:

- Using an Ordered Sequence
- Using Search Trees - binary search trees, AVL trees, red-black trees, (2,4)-trees

Implementing MAPs with an Unordered Sequence

- *unordered sequence*



- *get, remove and put takes $O(n)$ time*
- *The "insert" part takes $O(1)$ time, but we need first to search for key duplicate which takes $O(n)$.*

Implementing a Map an Ordered Sequence


- *array-based ordered sequence* (assumes keys can be ordered)



- searching takes $O(\log n)$ time (binary search)
- inserting and removing takes $O(n)$ time
- application to look-up tables (frequent searches, rare insertions and removals)

Binary Search

- narrow down the search range in stages
- “high-low” game
- Example: `get(7)`





if low > high then

```
else    mid  $\leftarrow$  (low+high) / 2
```

```
return key(mid)
```

```
return BinarySearch(S, k, low, mid-1)
```

2	4	5	7	8	9	12	14	17	19	22	25	27	28	33	37
↑							↑								↑
low							mid								high



Running Time of Binary Search

- The range of candidate items to be searched is *halved after each comparison*

comparison	search range
0	n
1	$n/2$
2	$n/4$
...	...
2^i	$n/2^i$
$\log_2 n$	1

In the array-based implementation, access by rank takes $O(1)$ time, thus *binary search runs in $O(\log n)$ time*


Binary Search Tree

Searching

Cost of Searching


Insertion


Deletion




Binary Search Trees

- A binary search tree is a binary tree T such that
 - each internal node p stores an item (k, v) of a MAP.
 - keys stored at nodes in the left subtree of p are less than k .
 - keys stored at nodes in the right subtree of p are greater than k .
 - external nodes do not hold elements but serve as place holders (dummy leaves).







Question: How can we traverse the tree so that we visit the elements in increasing key order?

IN-ORDER TRAVERSAL

always traverses the keys in increasing order
in a binary search tree !!!

MAP Operations using Binary Search Trees

Searching `get(k):`
 use `TreeSearch(k)`

Inserting/updating value `put(k, v):`
 use `TreeInsert(k,v)`

Removing `remove(k):`
 use `TreeDelete(k)`

Search

- To search for a key k , we trace a downward path starting at the root
- The next node visited depends on the outcome of the comparison of k with the key of the current node
- If we reach a leaf, the key is not found and we return this dummy leaf which will help with insertion if needed.
- Example: `TreeSearch(4)`

Algorithm `TreeSearch(k)`

`TreeSearch(root,k)`

Procedure `TreeSearch(p, k)`

if p is external then

return p {unsuccessful search}

else if $k == \text{key}(p)$


return p {successful search}

else if $k < \text{key}(p)$

return `TreeSearch(left(p),k)`


else { $k > \text{key}(v)$ }

return `TreeSearch(right(p),k)`



Search Example I

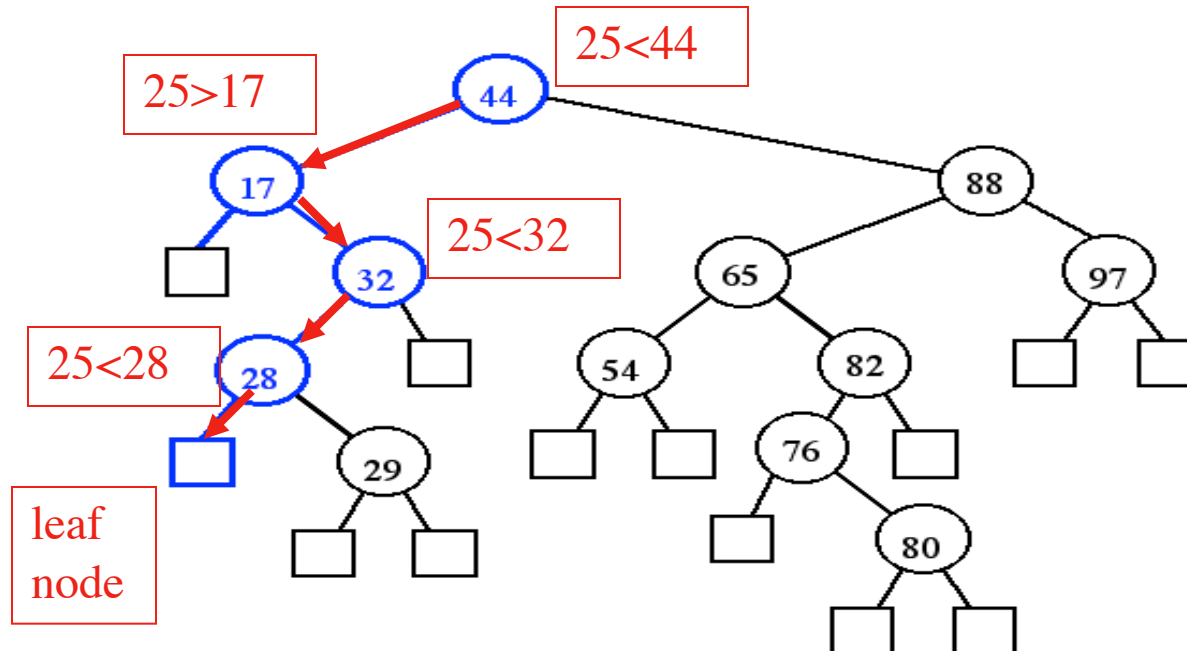
Successful **TreeSearch**(76)



- A successful search traverses a path starting at the root and ending at an internal node.


Search Example II

Unsuccessful **TreeSearch**(25)




- An unsuccessful search traverses a path starting at the root and ending at an external node

Cost of Search

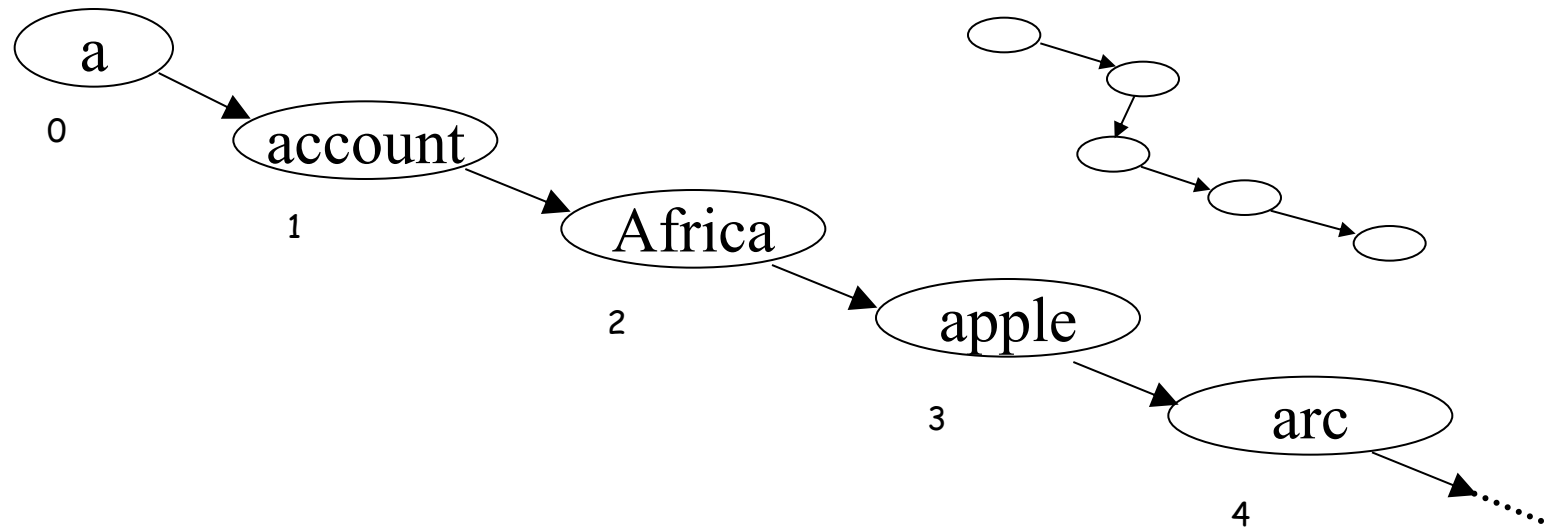


Worst tree



Best tree

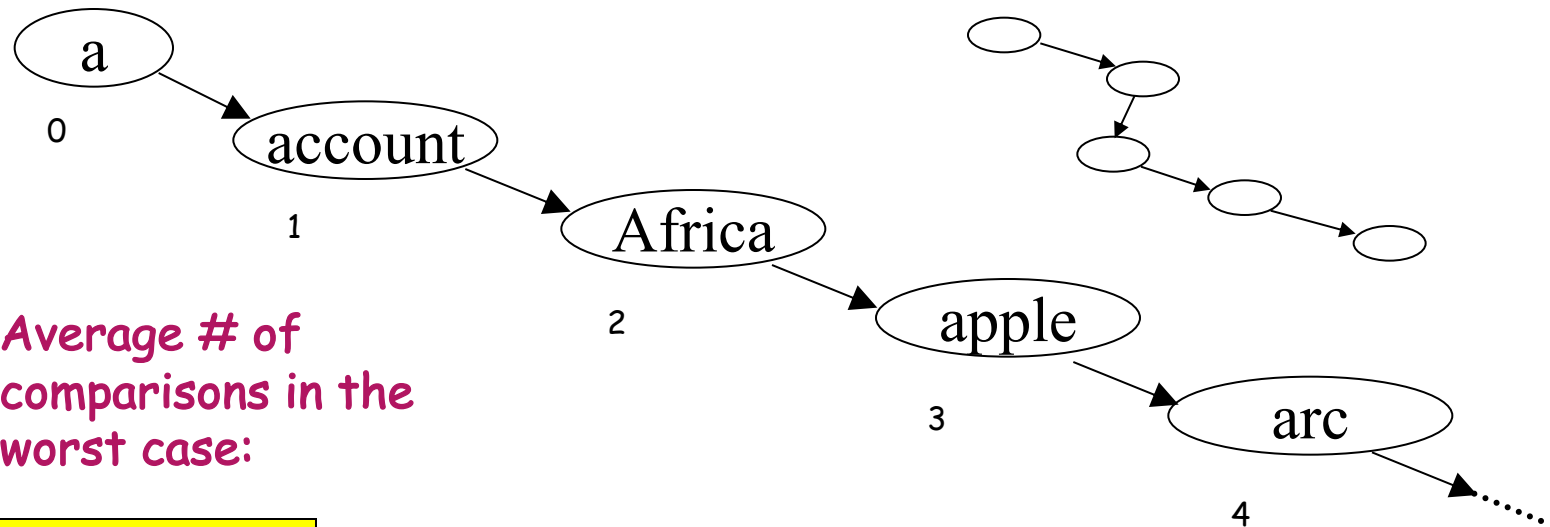
Cost of Search: Worst Case



Worst possible Tree:
Worst Case:

$O(n)$

Cost of Search - Average Worst Case

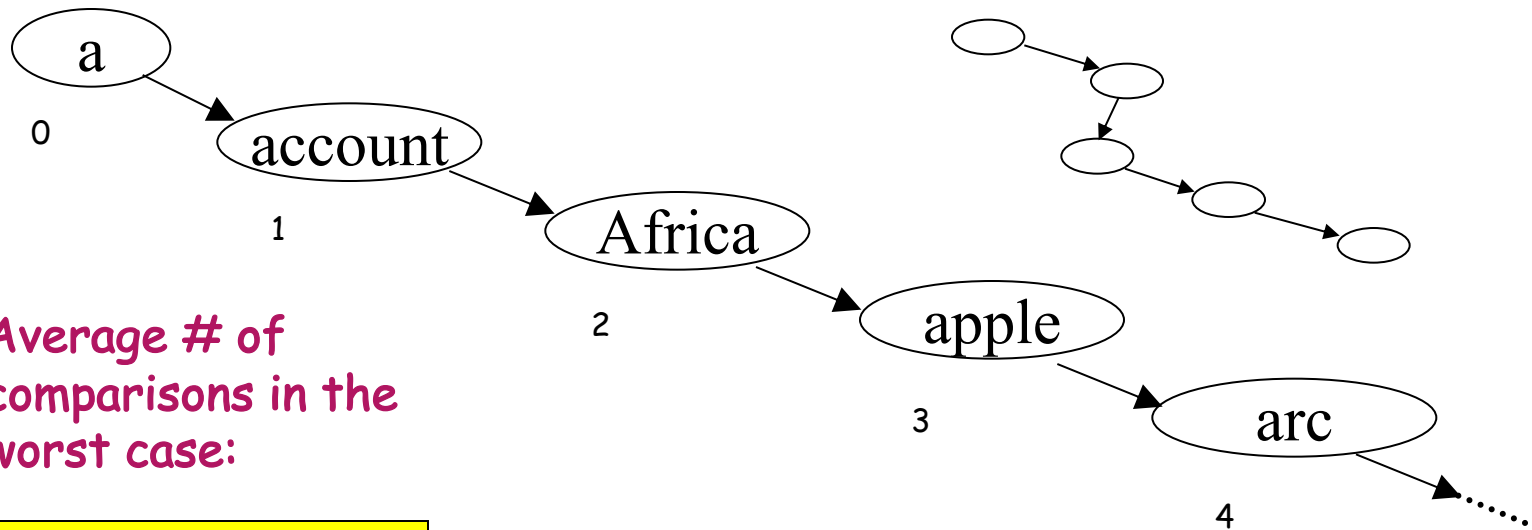


Successful
search

Path to node i has length i , to get there we do i comparisons

$$\text{Avg cost} = (1/n) \sum i = O(n)$$

Cost of Search - Average Worst Case






Average # of
comparisons in the
worst case:

Unsuccessful
search

An unsuccessful search
always takes $O(n)$ comparisons
for n internal nodes

Cost of Search: Best Case






Leaves are on the same level or on an adjacent level.

Length of path from root to node $i = \lfloor \log i \rfloor$

Worst case of the best possible tree: $O(\log n)$

Cost of Search: Average Best Case



Leaves are on the same level or on an adjacent level.

Length of path from root to node $i = \lfloor \log i \rfloor$


Comparisons to node i : $\log i$

→ Average # of comparisons in the best possible tree

Successful
search

$$\frac{1}{n} \sum_{i=1}^n \log i = O((n \log n) / n) = O(\log n)$$

Comparisons to node i : path of length (integer part of) $\log i$






Ex: to node 2: $\log 2 = 1$

to node 4,5,6,7:
 $\log 4 = \log 5 =$
 $\log 6 = \log 7 = 2$

Comparisons to node i : $O(\log i)$

Cost of Search: Average Best Case



Leaves are on the same level or on an adjacent level.

Length of path from root to node $i = \lfloor \log i \rfloor$

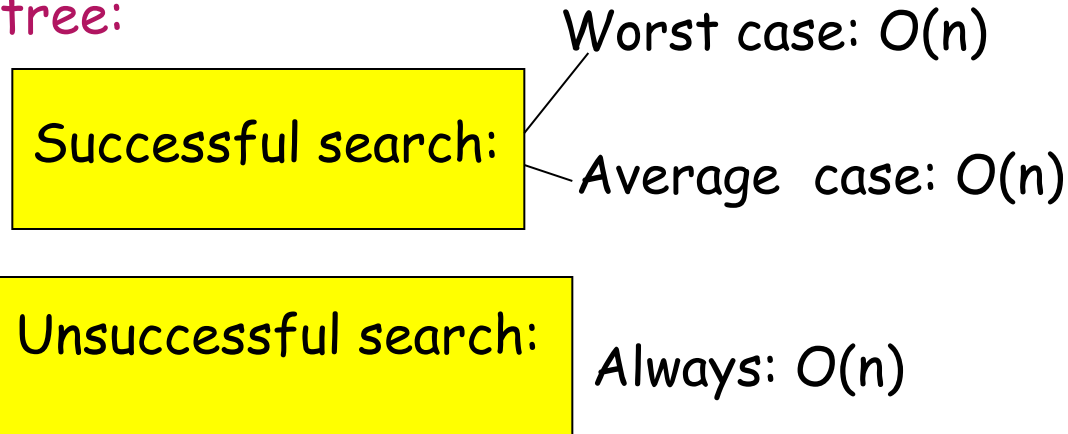
Only paths to external nodes count.

Unsuccessful
search

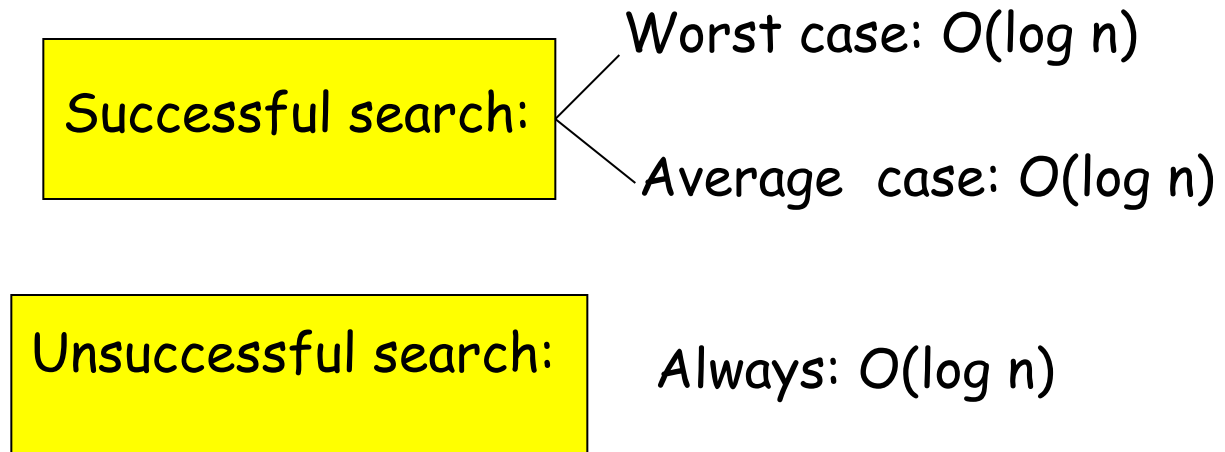
always $O(\log n)$

Summary

Worst tree:




Best Tree:



Note: worst case search for arbitrary binary search tree $O(n)$ 29


Insertion case I

- To perform `TreeInsert(k, v)`, let `w` be the node returned by `TreeSearch(k, T.root())`
- If `w` is external, we know that `k` is not stored in `T`. We call `expandExternal(w, (k,v))` to store `(k, e)` in `w`



expandExternal(p,(k,v)):

Transform p from an external node into an internal node by creating two new children



expandExternal(p,(k,v)):

if isExternal(p)

create new nodes new1 and new 2

p.left \leftarrow new1

p.right \leftarrow new2

store entry (k,v) in p

size \leftarrow size +2

Insertion case II

- If w is internal, we know the item with key k is stored at w .
In this case, we just replace the value on this node to the given value v .

Insertion in a Binary Search Tree

```
Algorithm TreeInsert(k,v)
    p = TreeSearch(root(),k)
    if k == key(p) then
        change p's value to (v)
    else
        expandExternal(p,(k,v)):
```


Construct a Tree



What would be the result of constructing a tree from repeated insertions of the following sequences?

- a. 5,8,3,7,1,9,2,4,6
- b. 1,2,3,4,5,6,7,8,9
- c. 5,4,6,3,7,2,8,1,9




When do you think trees work best?



Deletion I

- To perform operation $\text{remove}(k)$, we search for key k
- Assume key k is in the tree, and let v be the node storing k
- If node v has a leaf child w , we remove v and w from the tree with operation $\text{removeAboveExternal}(w)$
- Example: remove 4



`removeAboveExternal(v):`







removeAboveExternal(v):

```

if isExternal(v) {
  p ← parent(v)
  s ← sibling(v)
  if isRoot(p) {
    s.parent ← null
    root ← s
  }
  else {
    g ← parent(p)
    if (p is leftChild(g)) g.left ← s
    else g.right ← s
    s.parent ← g
  }
  size ← size - 2
}

```







Deletion II


- We consider the case where the key k to be removed is stored at a node p whose children are both internal

- we find the internal node w that follows p in an inorder traversal (note w it does not have a left child!)
- we copy $entry(w)$ into node p
- we remove node w and its left child z (which must be a leaf) by means of operation $removeAboveExternal(z)$

- Example: $remove(3)$

Note: see textbook for different approach: locating node w that precedes p in inorder traversal. How would this change the method above?





Practice, practice, practice...

- a. Delete the 3 from the tree you got in exercise (a) in page 33.
- b. Now delete node 5.


Cost of Inserting and Deleting = Cost of Search

Summary:

Consider a dictionary with n items implemented by means of a binary search tree of height h

- the space used is $O(n)$
- methods `findElement`, `insertItem` and `removeElement` take $O(h)$ time


The height h is $O(n)$ in the worst case and $O(\log n)$ in the best case




Conclusion

- To achieve good running time, we need to keep the tree *balanced*, i.e., with $O(\log n)$ height.
 - Various balancing schemes can be explored:
 - AVL trees* and *red-black trees* are a balanced binary search trees: their height is $O(\log n)$
 - A *(2,4)-tree* is a search tree (not binary, each internal node has 2, 3 or 4 kids); its height is also $O(\log n)$.
- Using simply a binary search tree gives worst case running time $O(n)$ for search, insert and delete operations!


Delete 13



Delete 13



Insert 16



Add 16

