


AVL Trees

Adel'son-Vel'skii and Landis

Data structure that implements MAP ADT


- Height of an AVL Tree
- Insertion and restructuring
- Removal and restructuring
- Costs



AVL Tree

- AVL trees are balanced.
- An AVL Tree is a *binary search tree* such that for every internal node v of T , the *heights of the trees rooted at the children of v can differ by at most 1*.


An example of an AVL tree where the heights are shown next to the nodes:



Balancing Factor

$\text{height}(\text{right subtree}) - \text{height}(\text{left subtree})$

$\in \{-1, 0, 1\}$ for AVL tree



Height of an AVL Tree


Note: “longest” possible heap with n nodes.

Always $O(\log n)$

But cannot search efficiently

Note: “longest” possible binary tree with n nodes:


$O(n)$



Inefficient to search in the worst case.

We'll now see that the *height* of an AVL tree T storing n keys is $O(\log n)$.


Note: AVL tree with the highest possible number of internal nodes for a given height h :



To construct the “longest” possible AVL tree, we look for the *minimum number of nodes of an AVL tree of height h .* $n(h)$


Easy to see that $n(1) = 1$ and $n(2) = 2$

(please note that dummy nodes are not shown but contributing to the height)



1

$$n(1) = 1$$




2

$$n(2) = 2$$

$n(h)$: the *minimum number of internal nodes* of an AVL tree of height h .




For $n \geq 3$, an AVL tree of height h contains the root node, one AVL subtree of height $h-1$ and the other AVL subtree of height $h-2$.



$$n(h) = 1 + n(h-1) + n(h-2)$$

Height of an AVL Tree

$$n(h) = 1 + n(h-1) + n(h-2)$$



Clearly:

$$n(h-1) > n(h-2)$$

so: $n(h) > 1 + n(h-2) + n(h-2)$
 $n(h) > 1 + 2 n(h-2) > 2 n(h-2)$

$$n(h) > 2n(h-2)$$

for any h

Height of an AVL Tree

So, now we know: $n(h) > 2$ $n(h-2)$ }
but then also: $n(h-2) > 2$ $n(h-4)$ } $n(h) > 4 n(h-4)$

$n(h) > 4n(h-4)$ }
but then also: $n(h-4) > 2n(h-6)$ } $n(h) > 8n(h-6)$

We can continue:

$$n(h) > 2n(h-2)$$

$$n(h) > 4n(h-4)$$

$$n(h) > 8n(h-6)$$

...

$$n(h) > 2^i n(h-2i)$$

$$n(h) > 2^{i(n(h-2i))}$$



$$h-2i = 2$$

And we know that

$$n(1) = 1$$

$$n(2) = 2$$

Now we pick i such that h is either 1 or 2

That is pick $i = \text{ceil}(h/2) - 1$, and substitute:

$$\text{for } i = \text{ceil}(h/2) - 1 \quad n(h) > 2^{\text{ceil}(h/2) - 1} n(1)$$

$$n(h) > 2^{h/2-1}$$

$$\log n(h) > \log 2^{(h/2)-1}$$

$$\log n(h) > h/2 - 1$$

$$h < 2 \log n(h) + 2 < 2 \log n + 2$$


which means that h is $O(\log n)$

Insertion


- A binary search tree T is called *balanced* if for every node v , the height of v 's children differ by at most one.
- Inserting a node into an AVL tree involves performing an *expandExternal(w)* on T , which changes the heights of some of the nodes in T .
- If an insertion causes T to become *unbalanced* we have to rebalance...

Insertion








Before



After left
insertion



After right
insertion





Rebalancing after insertion

We are going to identify 3 nodes which form a grandparent, parent, child triplet and the 4 subtrees attached to them. We will rearrange these elements to create a new balanced tree.

Rebalancing

Step 1: Trace the path back from the point of insertion to the first node whose **grandparent is unbalanced**. Label this node x, its parent y, and grandparent z.



Examples



Rebalancing

Step 1: Trace the path back from the point of insertion to the first node whose **grandparent is unbalanced**. Label this node **x**, its parent **y**, and grandparent **z**.



Examples



Rebalancing

Step 2: These nodes will have 4 subtrees connected to them. Label them T_1, T_2, T_3, T_4 from left to right.


Examples



Rebalancing


Step 3: Rename x, y, z to a, b, c according to their inorder traversal i.e. if y, x, z is the relative order of those nodes following the inorder traversal then label y 'a', x 'b' and z 'c'.

Example





Rebalancing

Step 4: Replace the tree rooted at z with the following tree:



Rebalance done!

Example: after inserting 54



Does this really work?

We need to see that the new tree is :



- a) A Binary search tree - the inorder traversal of our new tree should be the same as that of the old tree

Inorder traversal: by definition is T1 a T2 b T3 c T4

- b) Balanced: have we fixed the problem?


We consider 2 types of examples

Example 1



Inorder: T_1 z T_2 y T_3 x T_4

Example 2



Inorder: $T_1 \ y \ T_2 \ x \ T_3 \ z \ T_4$

An Observation...

Notice that in both cases, the new tree rooted at b has the same height that the old tree rooted at z had before insertion.

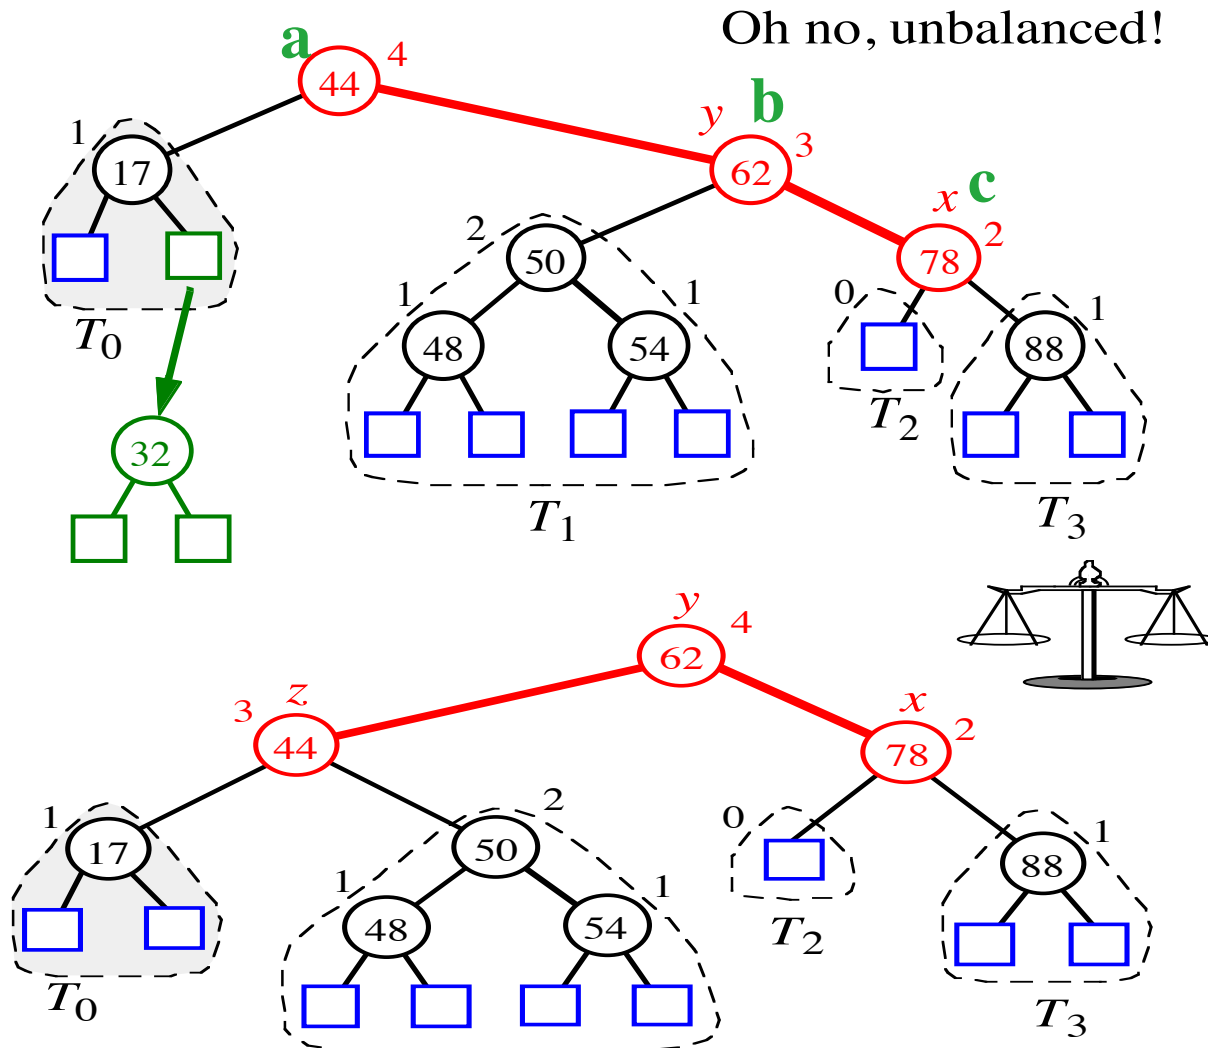
So.. once we have done one rebalancing act, we are done.

Removal

- We can easily see that performing a `removeAboveExternal(w)` can cause T to become unbalanced.
- Let z be the **first unbalanced** node encountered while travelling up the tree from w . Also, let y be **the child of z with the larger height**, and let x be **the child of y with the larger height**.
- We can perform operation `restructure(x)` to restore balance at the subtree rooted at z .
- As this restructuring may upset the balance of another node higher in the tree, we must continue checking for balance until the root of T is reached

Removal (contd.)


the choice of x is not unique !!!




Whew,
balanced!

Removal (contd.)

- we could choose a different x :



Oh no,
unbalanced!



Whew,
balanced!

COMPLEXITY

Searching:	findElement(k):
Inserting:	insertItem(k, o):
Removing:	removeElement(k):

$O(\log n)$

Some implementation details are very important:

The trinode restructure is accomplished using the rotation operation:




Figure 11.8: A rotation operation in a binary search tree. A rotation can be performed to transform the left formation into the right, or the right formation into the left. Note that all keys in subtree T_1 have keys less than that of position x , all keys in subtree T_2 have keys that are between those of positions x and y , and all keys in subtree T_3 have keys that are greater than that of position y .

Trinode restructuring using rotation operation:

474

Chapter 11. Search Trees

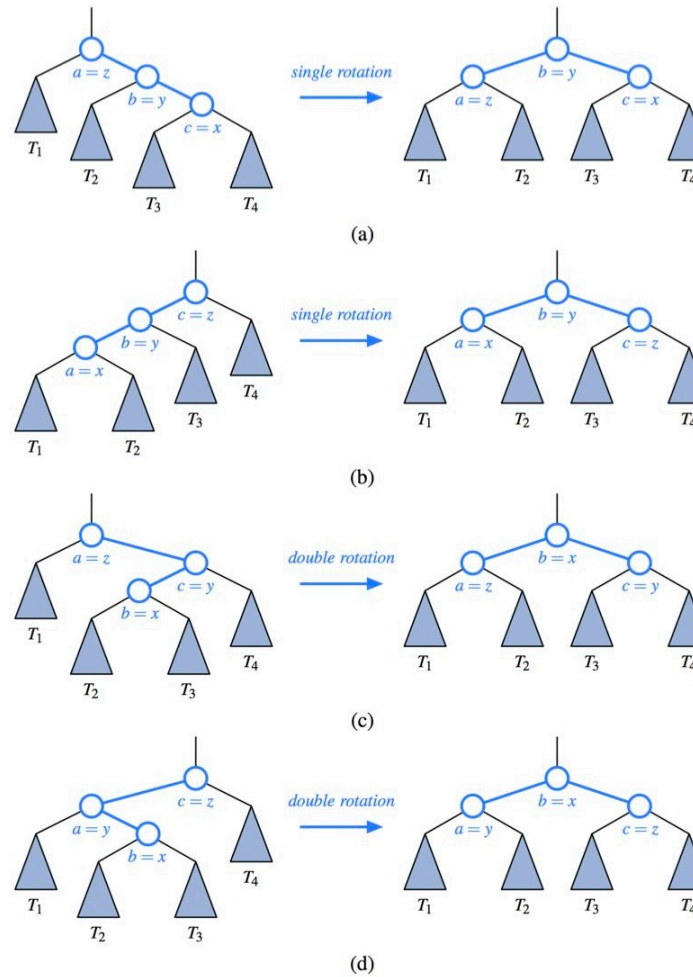


Figure 11.9: Schematic illustration of a trinode restructuring operation: (a and b) require a single rotation; (c and d) require a double rotation.

Rotate:

```
28  /** Relinks a parent node with its oriented child node. */
29  private void relink(Node<Entry<K,V>> parent, Node<Entry<K,V>> child,
30                  boolean makeLeftChild) {
31      child.setParent(parent);
32      if (makeLeftChild)
33          parent.setLeft(child);
34      else
35          parent.setRight(child);
36  }
37  /** Rotates Position p above its parent. */
38  public void rotate(Position<Entry<K,V>> p) {
39      Node<Entry<K,V>> x = validate(p);
40      Node<Entry<K,V>> y = x.getParent(); // we assume this exists
41      Node<Entry<K,V>> z = y.getParent(); // grandparent (possibly null)
42      if (z == null) {
43          root = x; // x becomes root of the tree
44          x.setParent(null);
45      } else
46          relink(z, x, y == z.getLeft()); // x becomes direct child of z
47      // now rotate x and y, including transfer of middle subtree
48      if (x == y.getLeft()) {
49          relink(y, x.getRight(), true); // x's right child becomes y's left
50          relink(x, y, false); // y becomes x's right child
51      } else {
52          relink(y, x.getLeft(), false); // x's left child becomes y's right
53          relink(x, y, true); // y becomes left child of x
54      }
55  }
56  /** Performs a trinode restructuring of Position x with its parent/grandparent. */
57  public Position<Entry<K,V>> restructure(Position<Entry<K,V>> x) {
58      Position<Entry<K,V>> y = parent(x);
59      Position<Entry<K,V>> z = parent(y);
60      if ((x == right(y)) == (y == right(z))) { // matching alignments
61          rotate(y); // single rotation (of y)
62          return y; // y is new subtree root
63      } else { // opposite alignments
64          rotate(x); // double rotation (of x)
65          rotate(x);
66          return x; // x is new subtree root
67      }
68  }
69  }
```

Code Fragment 11.10: The `BalanceableBinaryTree` class, which is nested within the `TreeMap` class definition (continued from Code Fragment 11.9).

Rebalancing operation for AVL insertions and deletions:

```
28  /**
29   * Utility used to rebalance after an insert or removal operation. This traverses the
30   * path upward from p, performing a trinode restructuring when imbalance is found,
31   * continuing until balance is restored.
32   */
33  protected void rebalance(Position<Entry<K,V>> p) {
34      int oldHeight, newHeight;
35      do {
36          oldHeight = height(p);                // not yet recalculated if internal
37          if (!isBalanced(p)) {                  // imbalance detected
38              // perform trinode restructuring, setting p to resulting root,
39              // and recompute new local heights after the restructuring
40              p = restructure(tallerChild(tallerChild(p)));
41              recomputeHeight(left(p));
42              recomputeHeight(right(p));
43          }
44          recomputeHeight(p);
45          newHeight = height(p);
46          p = parent(p);
47      } while (oldHeight != newHeight && p != null);
48  }
49  /** Overrides the TreeMap rebalancing hook that is called after an insertion. */
50  protected void rebalanceInsert(Position<Entry<K,V>> p) {
51      rebalance(p);
52  }
53  /** Overrides the TreeMap rebalancing hook that is called after a deletion. */
54  protected void rebalanceDelete(Position<Entry<K,V>> p) {
55      if (!isRoot(p))
56          rebalance(parent(p));
57  }
58 }
```


At this point in the class I discuss how AVL trees are implemented in the 6th edition of the textbook by Goodrich, Tamassia and Goldwasser.

Please, refer to pages:

466-470 class `TreeMap<K,V>`

Note methods: `put(K key, V value)`, `remove(K key)`

475-478 class `BalancedBinaryTree<K,V>`

Note: hooks for rebalancing present in `TreeMap`,

Methods: `rotate`, `restructure`

486-487 class `AVLTreeMap<K,V>`

Note methods: `rebalanceInsert`, `rebalanceDelete`,
`rebalance`